

From Scalar Compute to Vector Compute

Song Liu (song.liu@bristol.ac.uk)

GA 18, Fry Building,

Microsoft Teams (search "song liu").

Scalar Compute in R is SLOW

- R is very similar to C when it comes to scalar computing (the operand is a single datum).
 - Similar syntax, with minor differences.
- However, we have seen in the previous lab that **scalar compute is very slow in R**.
 - Particularly so when **loops** are involved.
- Is R really slow then?

Matrix Multiplication in R

- `%*%` is matrix multiplication operator in R.
 - `A %*% B` computes matrix multiplication between matrix `A` and `B`,
- Let us multiply two 1000 by 1000 matrices in R:

```
# generating two 1000X1000 random matrices
A <- matrix(rnorm(1000*1000), nrow = 1000, byrow = T)
B <- matrix(rnorm(1000*1000), nrow = 1000, byrow = T)

# built-in matrix multiplication
start = Sys.time()
C <- A%*%B
end = Sys.time()

print(end - start)
```

Matrix Multiplication in R

- It takes R 1 second to do the above multiplication.
- It takes C 5.7 second to do the same multiplication.
 - Using the code we wrote in lab 7.
- Why does R outperform C in matrix multiplication?

Single Instruction Single Data

- When C was born (1970s), CPU designs are very primitive.
- CPU can only process **a handful numbers at a time**.
 - One **single** CPU instruction can only work on a **single** data set.
 - For example, an `add` instruction can only add two numbers at a time.
 - This processing paradigm is called **Single Instruction Single Data (SISD)**.
- No doubt, SISD is inefficient in data science programs with large data sets.

Single Instruction Multiple Data

- However, modern CPUs (starting from 1990s) support a feature called Single Instruction Multiple Data (SIMD).
 - One single CPU instruction can operate on a large number of data sets.
 - For example, a single SIMD `add` instruction can add multiple pairs of number at the same time.
 - Popular SIMD instruction sets include MMX, SSE, SSE2, AVX2, AVX512.
- SIMD makes CPUs suitable for vector compute.
- R (and many other language that appeared later than C) supports vector compute natively via SIMD.

Vector Compute

- When R interpreter executes `A*%B`, it automatically **translates your code into SIMD instructions** then execute it on CPU efficiently.
- By default, C compiler will **not** translate your code into more efficient SIMD instructions.
- Thus, R is faster in the matrix multiplication example.

Vector Compute

- In fact, R is a kind of `array programming` language, that are designed for processing large vector/matrix.
- Accelerated vector compute makes R suitable for statistics, machine learning and many other scientific and engineering domains.

To Sum Up

- Scalar Compute: $C > R$.
- Vector Compute: $R > C$.
 - C can support vector compute via other extension libraries, such as Intel's Math Kernel Library (MKL), or NVIDIA's Compute Unified Device Architecture (CUDA).

Common Vector Operators

- You can find reference to most of these from ART 2.4
- Vector Addition adds each dimension of vectors.

```
a <- c(1, 2, 3, 4)
b <- c(1, 2, 3, 4)
c <- a+b
[1] 2 4 6 8
```

- Vector Subtraction `-`, Multiplication `*`, Division `/` works the same way.
 - Do not confuse `*` with `%*%`, `*` is dimension-wise multiplication.

Common Vector Operators

- Scalar + vector means adding each element in the vector by the scalar value.

- ```
a <- c(1, 2, 3, 4)
print(a+1)
[1] 2 3 4 5
```

- subtraction/multiplication/division works in the same way.

- Many math functions operates on each element of the input vector ( sin, cos, exp, log ):

- ```
exp(c(1, 2, 3, 4))
[1] 2.718282 7.389056 20.085537 54.598150
```

Indexing Element(s) in a Vector

- Use `[]` to index a single element
 - The index of the first element is 1 not 0!
 - ```
a <- c(1, 2, 3, 4)
print(a[1])
[1] 1
print(a[5])
[1] NA
```
  - Index the element that beyond the range of the vector will return `NA`.

# Indexing Element(s) in a Vector

- You can index more than one elements at a time, simply by using another vector as indices:

```
a <- c(1, 2, 3, 4)
b <- c(1, 3, 4)
print(a[b])
[1] 1 3 4
#or simply
print(a[c(1, 3, 4)])
[1] 1 3 4
```

- You can also use conditional expression to index a vector:

```
a <- c(1, 2, 3, 4)
print(a[a>2])
[1] 3 4
```

## : Symbol

- `:` symbol generates a vector of a range of values.

```
a <- 1:3
print(a)
[1] 1 2 3
```

- Use `:` to access multiple contiguous elements in a vector.

```
a <- c(5, 6, 7, 8)
print(a[1:3])
[1] 5 6 7
```

# for Loop

- Syntax `for (variable in vector)`
- Two for loops below are the same.

- ```
for (element in a){  
    print(element)  
}
```

- ```
for (i in 1:length(a)){
 print(a[i])
}
```

# Matrix Construction

- You can create a matrix from a vector using `matrix` function:

```
a <- c(1, 2, 3, 4)
A <- matrix(a, nrow = 2)
print(A)
```

|       | [, 1] | [, 2] |
|-------|-------|-------|
| [1, ] | 1     | 3     |
| [2, ] | 2     | 4     |

- `nrow` parameter specifies the number of rows in the matrix.

```
a <- c(1, 2, 3, 4)
A <- matrix(a, nrow = 1)
print(A)
```

|       | [, 1] | [, 2] | [, 3] | [, 4] |
|-------|-------|-------|-------|-------|
| [1, ] | 1     | 2     | 3     | 4     |



# Matrix Construction

- `matrix` will always fill out the matrix by columns. If you want to fill out the matrix by rows:

```
A <- matrix(c(1,2,3,4), nrow = 2, byrow = T)
print(A)
```

|       | [, 1] | [, 2] |
|-------|-------|-------|
| [1, ] | 1     | 2     |
| [2, ] | 3     | 4     |

# Matrix Index

- Elements in matrices can be indexed by `[,]`.

```
A <- matrix(1:4, nrow = 2)
print(A[1,1])
[1] 1
print(A[1:2,1])
[1] 1 2
```

# Matrix Operators

- `+, -, *, /` works on matrix in element-wise fashion too.

```
A <- matrix(c(1,2,3,4), nrow = 2, byrow = T)
B <- matrix(c(1,2,3,4), nrow = 2, byrow = T)
print(A+B)
```

|      | [,1] | [,2] |
|------|------|------|
| [1,] | 2    | 4    |
| [2,] | 6    | 8    |

- `%*%` works as matrix multiplication

```
print(A%*%B)
```

|      | [,1] | [,2] |
|------|------|------|
| [1,] | 7    | 10   |
| [2,] | 15   | 22   |

# Matrix Operators

- Matrix Transpose is done by `t` function.

```
a <- matrix(1:6, nrow = 2)
print(a)
```

|      | [,1] | [,2] | [,3] |
|------|------|------|------|
| [1,] | 1    | 3    | 5    |
| [2,] | 2    | 4    | 6    |

```
t(a)
```

|      | [,1] | [,2] |
|------|------|------|
| [1,] | 1    | 2    |
| [2,] | 3    | 4    |
| [3,] | 5    | 6    |

# Matrix Operators

- Matrix Inversion: `solve` function.

```
A <- matrix(1:4, nrow = 2)
print(A%*%solve(A))
```

|       | [, 1] | [, 2] |
|-------|-------|-------|
| [1, ] | 1     | 0     |
| [2, ] | 0     | 1     |

# Conclusion

- R uses **SIMD instructions to accelerate vector compute** while C does not by default.
- Scalar Compute: C is faster than R.
- Vector Compute: R is faster than C.
- Most operations (such as `+, -, *, /, sin, exp`) in R are applied on the whole vector **at the same time**.
  - Hence they are accelerated by SIMD and are fast.

# Homework

1. Try commands