

Tutorial: Rejection Sampling

Song Liu (song.liu@bristol.ac.uk)

GA 18, Fry Building,

Microsoft Teams (search "song liu").

Sampling

- Sampling is one of the most common tasks in statistical data analysis and in particular Monte Carlo simulations.
- The task of sampling is to draw samples of a random variable, given a probability density/mass function.

Sampling

- For example, how do you generate samples from a normal distribution $N(\mu, \sigma)$ given its probability density function $p(\mathbf{x}) \propto \exp(-(x - \mu)^2 / 2\sigma^2)$?
- Of course, most data science programming languages have dedicated built-in function to generate samples from normal distributions.
 - ```
a <- rnorm(1000, mean = 1, std = .5)
```
  - Many basic IoT devices does not have full support of a wide variety of statistical functions.

# Rejection Sampling

- Suppose we have a uniform sampler that samples observations from  $U(0, 1)$  and a basic sampler for a known distribution with density function  $q(x)$
- Target: sample  $n_{\max}$  observations from a distribution with density function  $p(x) \propto \bar{p}(x)$ .
- Algorithm:
  - Draw a sample  $x$  from  $q(x)$  and a uniform sample  $u \sim U(0, 1)$
  - If  $u < \frac{p(x)}{Mq(x)}$ , accept sample  $x$ .
    - If # accepted samples equals to  $n_{\max}$ , quit.
  - Repeat the algorithm.

# Rejection Sampling

- $M$  in the previous algorithm is a big number, which needs to satisfy  $\frac{p(x)}{q(x)} \leq M$ .
  - In our toy experiment, we can set  $M$  to be a big number, say 100 or 1000.
- When "accepting" a sample, we save it to a vector, so the vector later contains only accepted samples.

# Vector Appending

- If `a` is a length- $k$  vector, `a[k + 1] = c` will append another element `c` at the end of `a`.

```
a <- c(1,2,3,4)
a[5] <- 5
print(a)
[1] 1 2 3 4 5
```

# Code Skeleton

```
how many samples do we want?
n_max = 1000

pbar <- function(x){
 # the PDF of distribution, from which you want to sample.
 return(exp(-(x-1)^2/.5))
}

acc <- c() # create an empty vector
n <- 0 # how many samples have we already obtained?
M <- 200 # M = 200
while(n < n_max){
 u <- runif(1) # generate a uniform sample
 x <- rnorm(1) # in this example, q(x) is N(0,1)

 # TODO: complete the algorithm here

}

print(mean(acc))
print(sd(acc))
```

# Questions:

1. Complete the above code skeleton according to the rejection sampling algorithm.
2. What output is expected when the algorithm is implemented correctly?
3. Fill out the blank: The bigger  $M$  is, the \_\_\_\_ (faster/slower) the algorithm becomes.